



BAND OF AGENTS HACKATHON 2026 · AI/ML API + FEATHERLESS AI · Brand.ai JUNE 18, 2026

VOICEHIRE

The competency exploration engine that interviews like a senior panel.

Most AI interviewers ask from a question bank. **VoiceHire has no questions —**
every probe is generated live, from what the candidate just said.

THE PROBLEM

Hiring a senior engineer costs 240 hours Not to strategy but to systemic chaos.

Existing tools optimize interview throughput. Candidates and hiring managers care about what was actually demonstrated.



CV Overload

Hundreds of resumes, zero signal.

Recruiters scroll until they blur.



Scheduling Chaos.

Engineers' calendars are impossible to align.

Candidates ghost early.



Gut-Feel Decisions

Weeks of interviews, yet we still decide on vibes, not evidence.

Existing tools optimize interview throughput. Candidates and hiring managers care about what was actually demonstrated.

THE SOLUTION

The end-to-end AI recruiting platform for **evidence-based hiring.**



Smart Sourcing.

Automated pipeline management that surfaces the right talent before the first call, reducing screening time by 70%.



Voice Interviewing.

Our core engine that explores competency like a senior panel, generating every probe live from candidate responses.



Evidence-based Hiring.

Structured portfolios that replace gut-feel with defensible, cited candidate data for every hiring decision.

From Smart Sourcing → Voice Interviewing → Evidence-based Hiring.

THE JOURNEY

From Interview Tool to Full Recruiting Workspace

The Full Workspace

A central hub for the entire hiring lifecycle, from sourcing to analytics.

- + Candidate Management
- + Job Posting & Tracking
- + AI Semantic Matching

IMPACT

3x Faster Hiring

Drastic reduction in time-to-hire through automated workflows.

95%

Match Accuracy Score

40%

Reduction in Manual Toil

REASONING

The agents argue with each other.

Most agents do what you ask. The Integrity Skeptic pushes back. When technical evidence claims high confidence but the transcript lacks the vocabulary to justify it, the Skeptic fires a challenge — visibly.

01 Overconfidence challenge

02 Vocabulary-vs-claim mismatch

03 Confidence-adjustment events

04 Visible chain-of-thought trace

This isn't an LLM wrapper. It's an agent network that knows what it doesn't know.

● INTEGRITY_CHALLENGE · DEEPSEEK R1 · SESSION #14

Senior Backend Engineer — Redis caching probe

⚠ OVERCONFIDENCE FLAG

Technical evidence scored 0.86 on 'Redis caching', but the candidate named the tool without describing an invalidation strategy — confidence likely overstated.

🧠 THINKING TRACE EXPOSED

DeepSeek R1 <think> reasoning captured: 'Candidate said we used Redis — no mention of TTL, eviction, or stale-read handling. Lower confidence to 0.45.'

📊 COVERAGE ADJUSTMENT

Coverage: 'Caching Strategy' moves WEAK → UNEXPLORED-AT-DEPTH. Session Brain re-queues for deeper follow-up.

📄 COMMITTEE NOTE QUEUED

'Surface-level only — recommend targeted follow-up before COVERED status.'

Probes generated live, before the candidate finishes their answer.

Deepgram Nova-2 → Qwen3 32B (probe) → aura-2-thalia-en / aura-2-jupiter-en → next question spoken in under 2s

FILLER

'That's a great example...'
cancelled when
probe was
ready in 1.6s.

... Maya Chen · Senior Backend Engineer

● LIVE · 18:42

🔊 PROBE GENERATED · 0:38

"You mentioned Redis for caching — walk me through how you handled cache invalidation when the underlying data changed."

📄 EVIDENCE EXTRACTED · 1:14

Candidate describes TTL-based expiry plus a write-through pattern for hot keys. Llama 4 Scout tags 2 competencies, confidence 0.78. → Coverage: WEAK → COVERED

⚡ NEXT TARGET SELECTED · 1:52

"Let's switch gears — tell me about a time a production incident forced you to make a tradeoff under pressure."

LATENCY

1.6s

listen → spoken

ENTERPRISE ANALYTICS

Data-driven hiring with automated candidate insights.



Resume Intelligence

AI-powered structured extraction from PDF/DOCX. Instant profile creation with skill-weighted scoring.

REACT

PYTHON

AWS

SYSTEM DESIGN

REDIS



Performance Trends

Visual KPIs and status breakdowns. Track completion rates, pipeline health, and match accuracy in real-time.



Audit-Ready Reports

Downloadable PDF portfolios with full transcripts, competency scorecards, and deliberation notes.



candidate_report_v1.pdf

INTEGRITY & LIVE MONITORING

Real-time integrity enforcement and session control.



4-Level Enforcement

From passive 'Observation Only' to strict 'Lockdown' mode with automated termination.



Client-Side Shield

Fullscreen locking, tab tracking, and copy/paste prevention ensure session purity.



Integrity Skeptic

DeepSeek R1 monitors every response for behavioral anomalies and LLM signatures.

MISSION CONTROL DASHBOARD

LIVE SESSION

System Architecture	85%
API Design	60%
Concurrency	45%
Security	90%

14:02 - Session Brain: Generating probe...

14:03 - Voice Persona: Audio streaming...

14:05 - Integrity: No violations detected.


14:08 - Evidence: Tagging skill 'API Design'.

14:10 - Session Brain: Adaptive turn...

THE CLOSED LOOP

Every utterance updates the map. The map decides the next probe.

Probe → answer → evidence → coverage delta → next target · Deterministic Python, zero extra LLM calls.



🗺️ The brain — coverage map

API Design · MUST HAVE	COVERED 0.91
Caching Strategy · MUST HAVE	COVERED 0.78
Incident Response · MUST HAVE	WEAK 0.34
Distributed Systems · NICE TO HAVE	UNEXPLORED —

⚖️ Evidence portfolio — committee

DECISION: HIRE — 8 of 9 must-have competencies COVERED

Technical Advocate: 'API design and caching evidence strong — two independent confirmations each.'

Evidence Critic: 'Incident Response remains WEAK — flag for follow-up before final offer.'

Committee Chair: 'Recommend HIRE with a 30-min technical follow-up on incident handling.'

AI MATCHING

Precision matching with GPT-4o-mini semantic intelligence.



Instant Semantic Analysis

Deep profile parsing across thousands of candidates in milliseconds.



Evidence-based Relevance

Matches based on proven competencies, not just keyword density.

\$0.003

PER SEARCH

95%

RELEVANCE SCORE

Top Candidates
AI-matched for: Senior Backend Engineer

Ranked by AI match score — 10 candidates found

Felix O'Connor
#1 match **95%**
Strong Match

Strengths

- Proficient in Go
- Extensive experience with Kubernetes
- Strong background in microservices

Gaps

- Limited experience with RESTful API design

"Felix O'Connor has a strong match with the required skills in Go and Kubernetes, along with relevant experience in building reliable systems. His background as a Site Reliability Engineer aligns well with the responsibilities of a Senior Backend Engineer."

Schedule Interview

Hiro Tanaka
#2 match **90%**
Strong Match

Strengths

- Strong expertise in Go
- Experience with Kubernetes and microservices

Gaps

- Less experience in direct backend development compared to others

"Hector Martinez has extensive experience in backend engineering with a focus on Go and Kubernetes, making him a strong candidate for the role. His leadership experience also adds value."

INFRASTRUCTURE

A multi-model stack built for latency and reliability.

RUBRIC SYNTHESIS**DeepSeek V4 Pro**

Complex reasoning & extraction

INTEGRITY MONITOR**DeepSeek R1**

Reasoning-based skepticism

PROBE GENERATION**Qwen3 32B**

Low-latency adaptive dialogue

EVIDENCE EXTRACTION**gpt-4o-mini**

High-precision semantic tagging

**CORE ENGINE****HIRING COMMITTEE****Qwen3 235B A22B**

Deep thinking & verdict consensus

**STT Pipeline**
Deepgram Nova-2**Real-time Sync**
WebSocket Streaming**TTS Pipeline**
Deepgram Aura-2

BUSINESS VALUE

Built for teams that interview *every week*, not once a quarter.

TODAY

Recruiter screen

Panel loop

Take-home

Manual debrief

Four stages. Interviewers improvise. Decisions rest on gut feel and a 200-word Slack summary.

4–6 hrs

of senior-engineer time per candidate



WITH VOICEHIRE

One session.

A live, adaptive interview the candidate experiences as a real conversation — and a structured evidence portfolio the committee can defend.

\$' 3144

per full session

EFFICIENCY

90×

demo sessions per \$10 hackathon credit

ONBOARD

1 JD

Paste a job description. Rubric Synthesizer does the rest.

AUDITABILITY

100%

Every competency verdict backed by cited evidence IDs.

TRACTION

Proving the model with real-world scale.

CANDIDATE PIPELINE

433.

Candidates parsed and indexed into the Recruiter Workspace.

EFFICIENCY GAIN

<3{

More demo sessions per \$10 credit compared to manual screening.

ACTIVE JOBS

10

Live job postings managed through the end-to-end platform.

UNIT ECONOMICS

\$0.11

Cost per full adaptive interview session (STT + LLM + TTS).

MATCH ACCURACY

92%

AI semantic matching precision across diverse job categories.

AUDITABILITY

100%

Every competency verdict backed by cited evidence IDs.

Your Interview Coach.

*That Adapts
Every Turn.*

A multi-agent voice interview system six specialized agents, two model providers, deterministic coverage tracking, and a hiring committee that produces a defensible evidence portfolio.



voicehire.dev/demo



[Demo](#)



[GitHub](#)




AI/ML API

Featherless



AG

Deepgram



Thanks for watching.